

The Array API Standard in SciPy


Lucas Colley, EuroSciPy 2024

Thursday 29th August 2024, 11:30–11:50, Room 6,
Maritime University of Szczecin, Poland



These Slides

<https://github.com/lucascolley/euroscipy24-slides>



Why? Introduction to Interoperability



Alice

Why? Introduction to Interoperability



Alice

Why? Introduction to Interoperability



Alice

Why? Introduction to Interoperability




Alice



Why? Introduction to Interoperability



Alice



Bob

Why? Introduction to Interoperability



Alice



Bob

Why? Introduction to Interoperability




Alice



Why? Introduction to Interoperability

- Use CuPy
- `np` -> `cp`



Alice

```
>>> import scipy
>>> import numpy as np
>>> x = np.asarray(...)
...
>>> y = scipy.some_function(x)
...
```

Why? Introduction to Interoperability

- Use CuPy
- `np` -> `cp`



Alice

```
>>> import scipy
>>> import numpy as np
>>> x = np.asarray(...)
...
>>> y = scipy.some_function(x)
...
```

```
>>> import cupy as cp
>>> x = cp.asarray(...)
...
>>> y = scipy.some_function(x)
KeyError: 'ALIGNED is not defined
for cupy.ndarray.flags'
```

Why? Introduction to Interoperability

- Replace SciPy with torch-magic (*fictional*)
- Use PyTorch
- `cp -> torch`




Alice

```
>>> import torch
>>> import torch_magic
...
>>> torch_magic.some_function(...)
# works!
```

Why? Introduction to Interoperability

- Replace SciPy with torch-magic (*fictional*)
- Use PyTorch
- `cp` -> `torch`



Alice

```
>>> import torch
>>> import torch_magic
...
>>> torch_magic.some_function(...)
# works!

>>> torch.expand_dims(...)
AttributeError: module 'torch' has
no attribute 'expand_dims'

# called 'unsqueeze' instead!
```

Why? Introduction to Interoperability

- Why can't SciPy use CuPy arrays / work on GPU?
- Why don't array libraries have the same API?
- Why does {some library} only work with one type of array?

About Me

- SciPy Maintainer (the algorithm library, not the conference!)
- Computer Science & Philosophy Undergraduate, Christ Church, University of Oxford
- Previously, an intern at Quansight Labs
- Based in Newcastle upon Tyne, UK



Why? Interoperability - SciPy

- SciPy is entirely built around NumPy arrays
- “Support for distributed arrays and GPU arrays” has been on SciPy’s roadmap for over 5 years.
- But “NumPy provides an array implementation that’s in-memory, CPU-only and single-threaded”^[1].

[1] https://data-apis.org/array-api/latest/use_cases.html#use-case-sciPy

Why? Interoperability - The Scientific Python Ecosystem

‘Array (provider) Libraries’



‘Array Consumer Libraries’

Credit: Aaron Meurer, ‘Python Array API Standard’, SciPy 2023

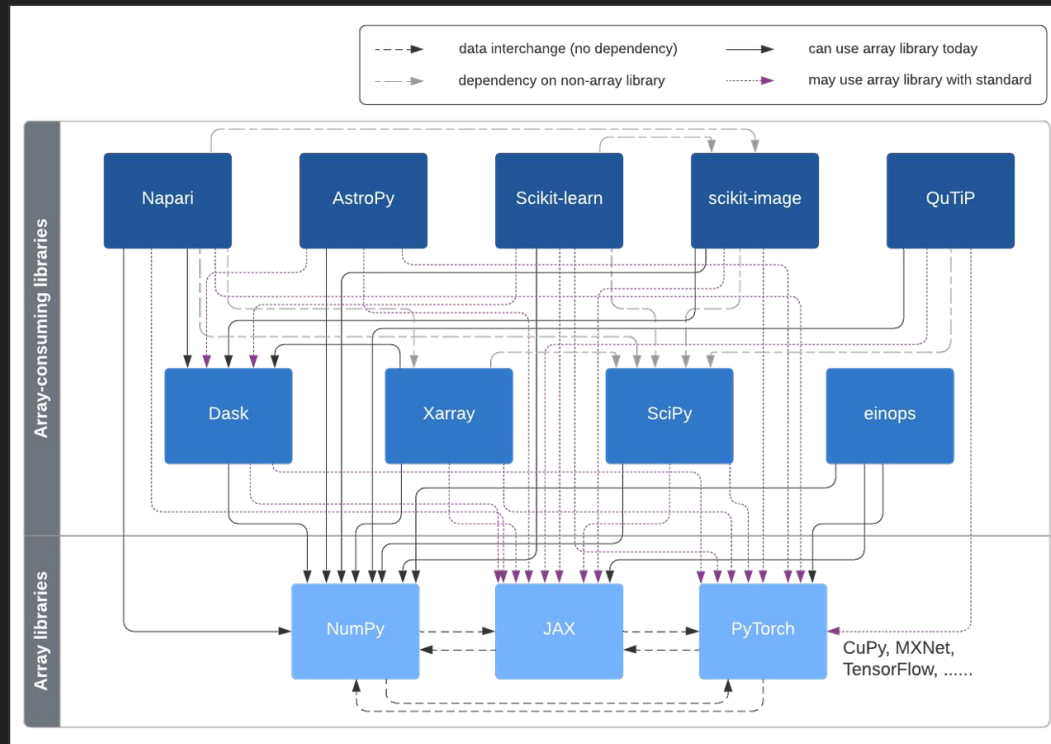
<https://github.com/data-apis/scipy-2023-presentation/blob/main/presentation/Slides.pdf>

Aside: It's Complex

A peek into the complex (actual and potential) interoperability within the Scientific Python Ecosystem (for illustrative purposes only):

Credit: Ivan Yashchuk & Ralf Gommers, 'A vision for extensibility to GPU & distributed support for SciPy, scikit-learn, scikit-image and beyond', 2021


<https://labs.quansight.org/blog/2021/11/pydata-extensibility-vision>



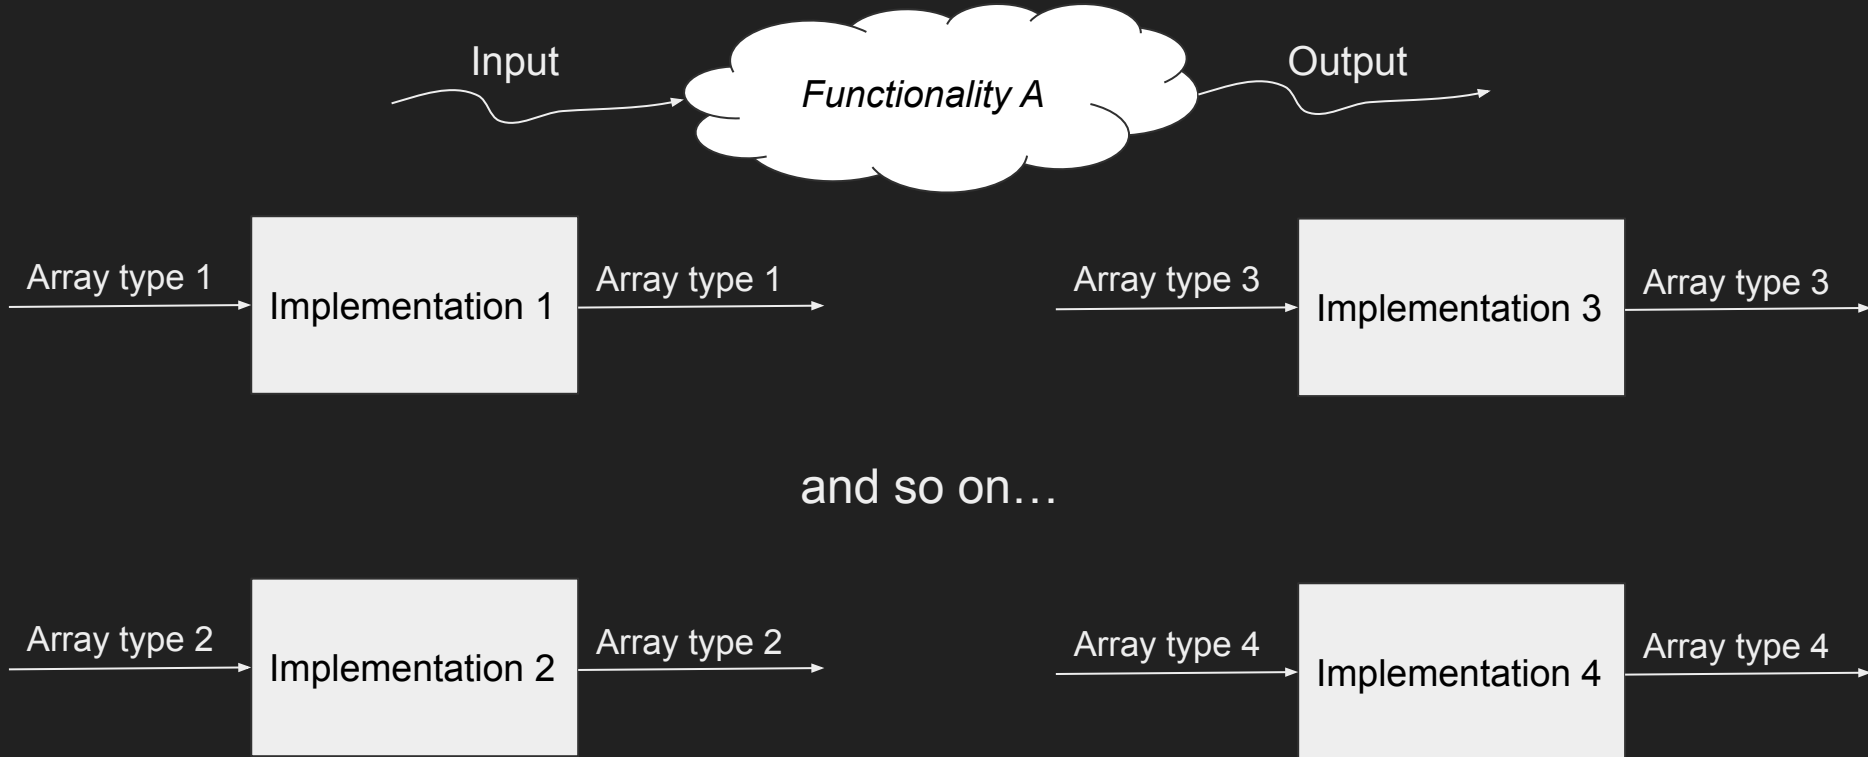
Why? Interoperability - The Bigger Picture

- Layered but divided ecosystem
- Interoperability is about much more than just bringing GPU support to the NumPy ecosystem

Why? Interoperability - The Current Ecosystem



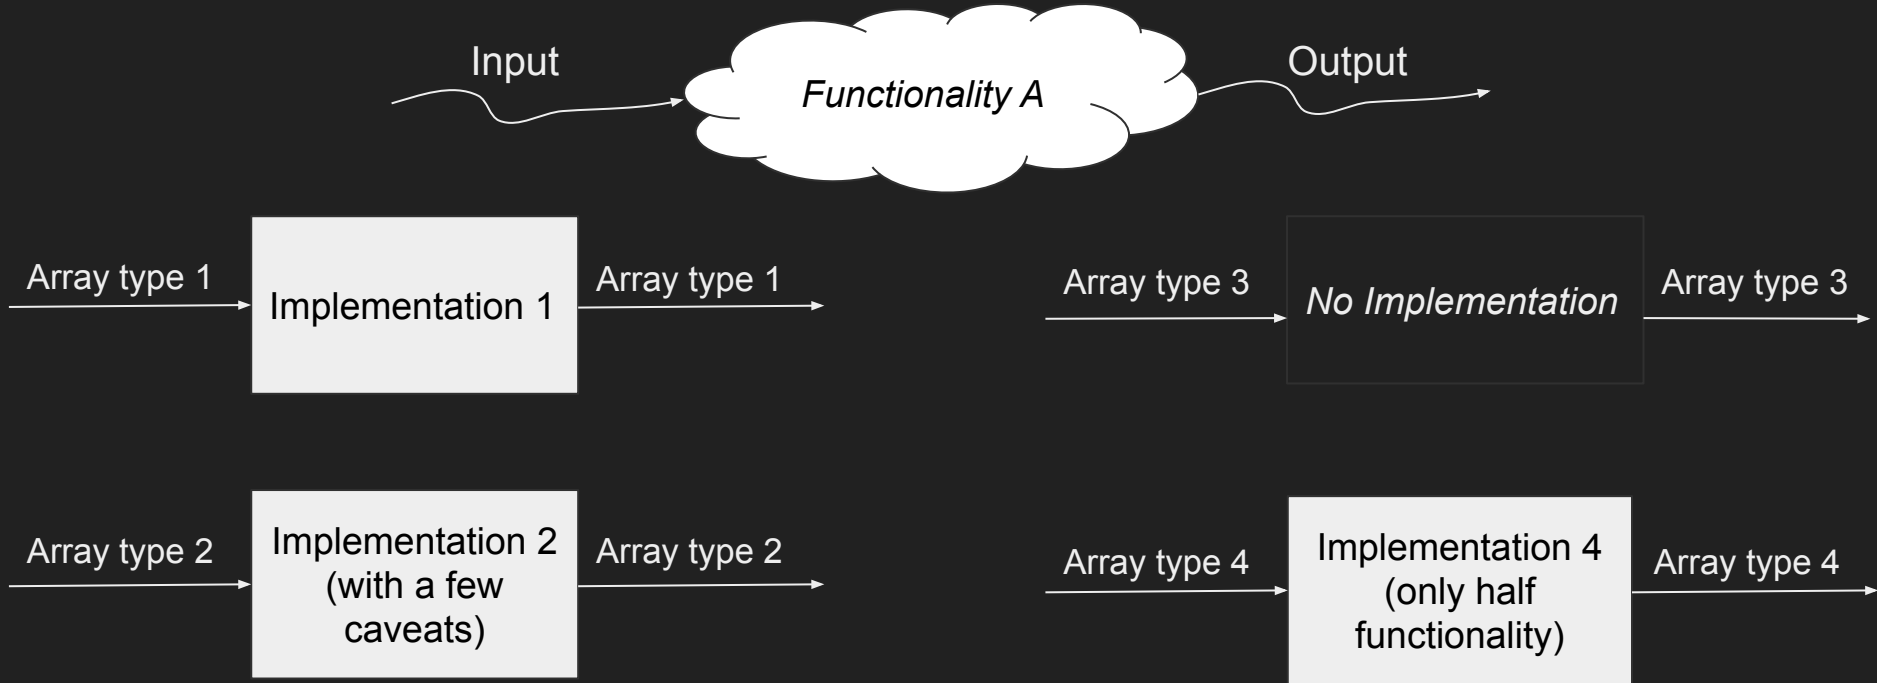
Why? Interoperability - The Current Ecosystem



Why? Interoperability - The Current Ecosystem

- Problem: maintenance cost!
- Difficult to make improvements in every implementation at once
- New array library? Needs another implementation
- All of this duplication does not seem efficient...

Why? Interoperability - The Current Ecosystem (really)



What? The (Python) Array API Standard

- Specifies a standard API for array libraries
- Starting from a minimal set of commonly implemented functionality
- Thus enabling “array-agnostic” implementations
- <https://github.com/data-apis/array-api>

Aside: The Python Array API Standard

- A lot more to say: past work, design principles, testing, methodology



<https://youtu.be/16rB-fosAWw>



Python Array API Standard

Toward Array Interoperability in the Scientific Python
Ecosystem

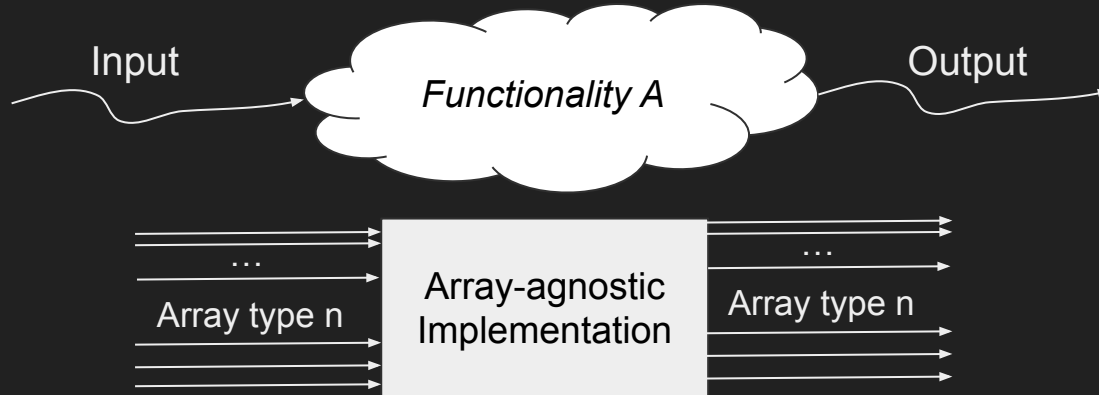


Aaron Meurer, Quansight
July 14, 2023
11:25–11:55, Amphitheater 204
SciPy 2023, Austin, TX



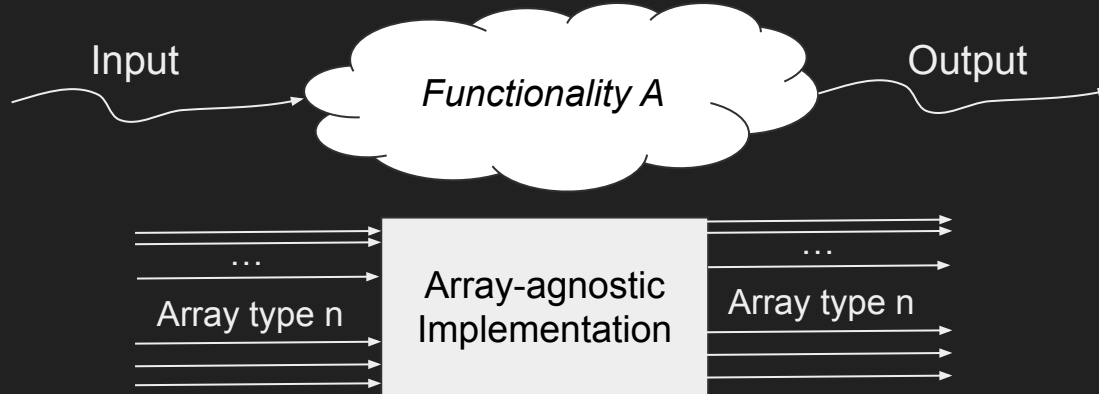
These slides
<https://github.com/data-apis/scipy-2023-presentation/blob/main/presentation/Slides.pdf>

What? Interoperability - A Unified Ecosystem



- Each array library implements the standard API
- Array-agnostic implementations which use just the standard API will work with all of them

What? Interoperability - A Unified Ecosystem



- Specialised implementations still possible (but only sometimes useful)



How? `xp = array_namespace(x)`

- The first line of most converted functions - get the namespace `xp` to use with the input arrays. Typically replace `np` with `xp`.

```
## scipy.cluster.vq.whiten
+ xp = array_namespace(obs)
+ obs = _asarray(obs, check_finite=check_finite, xp=xp)
- obs = _asarray_validated(obs, check_finite=check_finite)
+ std_dev = xp.std(obs, axis=0)
- std_dev = obs.std(axis=0)
  zero_std_mask = std_dev == 0
+ if xp.any(zero_std_mask):
- if zero_std_mask.any():
    std_dev[zero_std_mask] = 1.0
    warnings.warn(...)
return obs / std_dev
```

How? `x.__array_namespace__`

- `array_namespace` queries for the namespace via the dunder method
- No runtime imports! The array stores the namespace.
- `array_namespace` does other things for us:
 - Returns NumPy unless an env variable `SCIPY_ARRAY_API` is set
 - If it is set, reject some array types we do not want to support (e.g. `numpy.matrix`, arrays with esoteric dtypes, unknown objects which NumPy can't coerce with `np.asarray`)
 - Fallback to NumPy for Python 'array-likes'
- Simple to wrap and customise

How? array-api-compat

- Wraps each array library to (practically) full compliance
- Has an `array_namespace` helper which we wrap - so we can write code that works with array libraries today (despite missing the dunder method)
- Currently wraps NumPy, CuPy, PyTorch, JAX, Dask, ndonnx and Sparse
- <https://github.com/data-apis/array-api-compat>

How? Testing

- array-api-strict is a strict minimal implementation of the standard
- If our tests pass with array-api-strict arrays, they should pass with arrays from any compliant library
- We have pytest marks/fixtures to parametrize tests with array libraries & skip/xfail tests for certain backends/devices
- We have testing helpers which work across backends, such as `xp_assert_close` to replace `np.testing.assert_allclose`
- Check for correct output namespace is built-in to the helpers
- <https://github.com/data-apis/array-api-strict>

How? Testing

```
+ @skip_xp_backends(cpu_only=True)
+ @array_api_compatible
  @pytest.mark.parametrize("func", ['dct', 'dst', 'dctn', 'dstn'])
  @pytest.mark.parametrize("type", [1, 2, 3, 4])
  @pytest.mark.parametrize("norm", [None, 'backward', 'ortho', 'forward'])
- def test_fftpack_equivalience(func, type, norm):
+ def test_fftpack_equivalience(func, type, norm, xp):
    x = np.random.rand(8, 16)
+    fftpack_res = xp.asarray(getattr(fftpack, func)(x, type, norm=norm))
+    x = xp.asarray(x)
    fft_res = getattr(fft, func)(x, type, norm=norm)
-    fftpack_res = getattr(fftpack, func)(x, type, norm=norm)

-    assert_allclose(fft_res, fftpack_res)
+    xp_assert_close(fft_res, fftpack_res)
```

How? SciPy and Delegation

- Compiled code is out of scope for the Python array API standard
- But SciPy has a lot, due to being on the border of the array library and array-consumer library divide (only ~60% of the codebase is Python)
- For some of `scipy.fft` and `scipy.linalg`, we can use the array API standard extensions
- There has been work on delegating to CuPy/PyTorch/JAX separately from the standard, for `scipy.special`, `scipy.ndimage`, `scipy.fft`, `scipy.signal`
- In practice, lots of overlap with array API standard work: `array_namespace`, testing
- If there is no library to delegate to, must convert to NumPy and back

Aside: Delegation, Dispatching

- But this is a different approach to interoperability - relies on other implementations existing
- Most important for code which needs or has big gains from specialised implementations in a compiled language
- Generalising to a dispatching mechanism is a challenge -
c.f. Quansight-Labs/uarray, networkx, scientific-python/spatch

When and Where? SciPy - Current Progress

- Testing with PyTorch & JAX CPU in CI, CuPy also locally (GPU CI coming)
- Pure Python + NumPy code supports all libraries on all devices
- For compiled code, support for libraries with CPU execution
- Just a few rough edges - e.g. in-place operations, fancy indexing
- Modules currently covered: `cluster`, `constants`, `datasets`, `fft`, `io`, `ndimage`
- Partial coverage in `special`, `stats`
- Work in progress in various places

When and Where? SciPy - Looking Forward

- Still *a lot* of SciPy to convert!
- Support for `dask.array` in progress
- End goal: full coverage
- Then can think about raising warnings for array types which will see changed behaviour, and flipping the switch

When and Where? Open Source Contributors

- SciPy tracker
- Also scikit-learn (sprint!)
- A library you already contribute to!
- Feel free to direct maintainers to this talk

When and Where? Array Consumer Library Maintainers

- Your own library!
- Especially for lightweight, pure Python + array library code
- Can start experimenting now
- Reach out on the array-api-compatible repo - can vendor or make a dependency
- Can copy utilities SciPy and Scikit-learn are using
- Utilities which are generally useful will be upstreamed to array-api-compatible in the future

When and Where? Array Library Maintainers


- Implement array API standard support in your main namespace!
 - `numpy`, `cupy`, `jax.numpy`, `sparse`, `ndonnx`
- Or contribute a wrapper to array-api-compat
 - `dask.array`, `torch` available now
- Feedback on specific parts of the array API standard:
 - <https://github.com/data-apis/array-api>
- Higher level feedback:
 - <https://github.com/data-apis/consortium-feedback>

With thanks to:



Ralf Gommers

Chair, Consortium for
Python Data API Standards;
Chair, SciPy Steering Council



Aaron Meurer

Maintainer:
array-api, array-api-compat,
array-api-strict




Matt Haberland

SciPy Maintainer




Pamphile Roy

Original Author of SciPy's
array API support machinery;
SciPy Maintainer



Tyler Reddy

SciPy Release Manager




Jake Bowhay

SciPy Maintainer



Evgeni Burovski

SciPy Maintainer



h-vetinari

SciPy Maintainer

Questions?



These Slides



SciPy Tracker



Array API Repo

<https://github.com/lucascolley/euroscipy24-slides>

<https://github.com/scipy/scipy/issues/18867>

<https://github.com/data-apis/array-api>